




OptimalBinningWoE

CHEAT SHEET · SHEET 1 OF 2 · BIN, SCREEN, TRANSFORM

Optimal binning and Weight of Evidence for credit scoring and risk modelling: 37 C++ binning implementations behind one R interface, from a raw feature store to a deployed scorecard.

version 1.13.6

```
install.packages("OptimalBinningWoE")
library(OptimalBinningWoE)
```



What the transform measures

Binning replaces a raw column by its bins; WoE replaces each bin by the log-ratio of the event and non-event distributions it holds.

$$\text{WoE}_i = \ln(p_i / q_i) \quad \cdot \quad \text{IV} = \sum_i (p_i - q_i) \cdot \text{WoE}_i$$

p_i = share of the events that fall in bin i ; q_i = share of the non-events. A positive WoE is a riskier bin.

IV STRENGTH BANDS (SIDDIQI, 2006)

Unpredictive	Weak	Medium	Strong	Suspicious
< 0.02	0.02 – 0.10	0.10 – 0.30	0.30 – 0.50	≥ 0.50

An IV at or above **0.50** is far more often leakage than signal, which is why `obwoe_select()` rejects it by default. Medium is where the workhorses of a scorecard sit.

Fit

STEP 1

```
obwoe(data, target, feature = NULL,
       min_bins = 2, max_bins = 7,
       algorithm = "auto",
       control = control.obwoe())
```

data Data frame holding the target and the candidates.
target Column name. Binary 0/1, or multinomial 0,1,2,...
feature NULL bins every column but the target.
min_bins Bin-count envelope the optimiser must respect.
max_bins
algorithm One of the 28 names, or "auto": jedi for a binary target, jedi_mwoe for a multinomial one.
control See `control.obwoe()` on sheet 2.

WHAT COMES BACK: AN **OBWOE** OBJECT

```
m <- obwoe(german, target = "default", max_bins = 6)

m$summary      # feature, type, algorithm, n_bins, total_iv
m$results$page # bin, woe, iv, count, count_pos, count_neg,
               # cutpoints, converged, iterations
m$target_type  # "binary" or "multinomial"
```

LOOK AT IT

`summary(m)` Feature table, `sort_by = "iv"`, decreasing.
`print(m)` Compact fit report.
`plot(m, type=)` "iv" ranking bar chart · "woe" profile of one feature · "bins" counts and event rate. Pass `feature=`, `top_n = 15`.

Screen

STEP 2

Two criteria govern admission: Information Value strength and guaranteed rank ordering. Every candidate returns a row, so an automatic verdict stays reviewable.

```
sel <- obwoe_select(m,
  detail = "summary",      # or "full": one row per bin
  iv_min = 0.02, iv_max = 0.5,
  require_monotonic = "numeric", # "all" | "none"
  monotonicity = "weak", # or "strict"
  min_bins = 2, max_bins = Inf,
  min_bin_pct = 0, allow_degenerate = FALSE,
  top_n = NULL, sort_by = "iv", decreasing = TRUE)
```

COLUMNS THAT CARRY THE VERDICT

total_iv, iv_class IV and its Siddiqi band.
ks, gini, auc Discrimination of the binned score.
monotonic Event-rate ordering, plus `monotonic_direction`, `n_violations`,
monotonic_strict spearman.
n_degenerate_bins Bins with no events or no non-events.
quality Excellent · Good · Fair · Rejected.
selected The flag. `reason` and `reason_desc` say why.

REASON CODES (JOINED BY ";")

OK	IV_BELOW_MIN	IV_SUSPICIOUS	NOT_MONOTONIC	TOO_FEW_BINS
TOO_MANY_BINS	SMALL_BIN	DEGENERATE_BIN	NOT_IN_TOP_N	BINNING_ERROR

```
keep <- sel$feature[sel$selected]
```

Transform

STEP 3

```
obwoe_apply(data, obj,
  suffix_bin = "_bin", suffix_woe = "_woe",
  keep_original = TRUE, na_woe = 0)
```

Adds `<feature>_bin` and `<feature>_woe` per binned column. `na_woe` is the fallback for an unseen category or an unmodelled missing value. Numerical intervals are half-open on the right, `(a, b]`.

THE MANUAL PATH, ONE COLUMN AT A TIME

`ob_cutpoints_num()` Bin a numeric vector at cut points you supply.
`ob_cutpoints_cat()` Same for a categorical vector.
`ob_apply_woe_num()` Map a numeric vector through a fitted result; `missing_values = c(-999)`.
`ob_apply_woe_cat()` Categorical counterpart; `missing_values = c("NA", "Missing", "")`.

Evaluate

```
g <- obwoe_gains(m, feature = "duration")
g <- obwoe_gains(scored, target = "default",
  feature = "score_decile",
  use_column = "direct", n_groups = 10)
```

`use_column` "auto" · "bin" · "woe" · "direct".
`sort_by` "id" (the algorithm's own order, default) · "woe" · "event_rate" · "bin".
`n_groups` Cut a continuous score into that many quantile groups.

18 STATISTICS PER BIN

count	count_pct	pos_count	neg_count	pos_rate	neg_rate	pos_pct
neg_pct	odds	log_odds	woe	iv	cum_pos_pct	cum_neg_pct
ks	lift	capture_rate				

```
g$metrics # ks, gini, auc, total_iv, ks_bin
plot(g, type = "cumulative") # "ks" | "lift" | "woe_iv"
```

Related helpers: `obwoe_gains_score()` for the bin-level engine behind a fitted result, `obwoe_gains_variable()` for any binned data frame plus a grouping column.

Prepare the raw column

```
ob_preprocess(feature, target,
  num_miss_value = -999, char_miss_value = "N/A",
  outlier_method = "iqr", # "zscore" | "grubbs"
  outlier_process = FALSE,
  preprocess = "both",
  iqr_k = 1.5, zscore_threshold = 3,
  grubbs_alpha = 0.05)
```

Missing values become a value of their own rather than a dropped row, so the bin that holds them carries its own WoE and stays visible in the model document.

```
ob_check_distincts(x, target) # cardinality and
                              # separation warnings
```

Runnable from a clean session

The Statlog (German Credit) benchmark ships with the package.

```
german <- read.csv(gzfile(system.file(
  "extdata", "germancredit.csv.gz",
  package = "OptimalBinningWoE")),
  stringsAsFactors = FALSE)

german$default <- 1L - german$credit_risk
german$credit_risk <- NULL

m <- obwoe(german, "default", max_bins = 6)
sel <- obwoe_select(m)
out <- obwoe_apply(german, m)
```



Choose, score, deploy

CHEAT SHEET · SHEET 2 OF 2 · ALGORITHMS, SCORECARD, SQL

28 algorithm names cover 37 implementations — 21 accept a numerical feature, 16 a categorical one, 9 both. Every one is callable on its own, and the whole pipeline is callable as a single function.

C++ via Rcpp / RcppEigen

obwoe_algorithms()
vignette("introduction")

The 28 algorithms

NUMERICAL · CATEGORICAL

FAMILY	NAME	N C	OPTIMISES
Information-theoretic	jedi	N C	Joint entropy-driven intervals (<i>default</i>)
	jedi_mwoe	N C	JEDI for a multinomial target
	mdl1p	N	Fayyad–Irani MDL stopping rule
	fast_md1p	N	MDLP with a monotonicity constraint
	dmiv	N C	Divergence measures (Zeng, 2013)
	ivb	C	IV by dynamic programming
Statistical merging	cm	N C	Enhanced ChiMerge, χ^2 on neighbours
	fetb	N C	Fisher's exact test on neighbours
	mob	N C	Monotonic optimal binning
Shape-constrained	ir	N	Isotonic regression (PAVA)
	m1blp	N	Monotonic risk + likelihood-ratio pre-bins
	mb1p	N	Monotonic binning by linear programming
	osl1p	N	Optimal supervised partitioning
	ldb	N	Local density binning
	lpdb	N	Local polynomial density binning
	gmb	C	Greedy merge under monotonicity
Exact optimisation	dp	N C	Dynamic programming, global optimum
	bb	N	Branch and bound
	m1lp	C	Mixed-integer formulation
	sb1p	C	Sequential bounded partitioning
Search & metaheuristic	udt	N C	Entropy-based tree partitioning
	sab	C	Simulated annealing
	mba	C	Agglomerative monotonic merging
	swb	C	Sliding window over ordered levels
Unsupervised & streaming	sketch	N C	Streaming quantile sketch, sketch_k
	ewb	N	Equal width, then IV refinement
	kmb	N	K-means initialisation
	ubsd	N	Standard-deviation cut points

Picking one

- **Must the WoE be monotone?** Numerical: `ir`, `m1blp`, `mb1p`, `mob`. Categorical: `gmb`, `mob`, `mba`.
- **Millions of rows, or a stream?** `sketch`, then `ewb` or `kmb`.
- **Need the provable optimum?** `dp` and `bb` for numerical, `m1lp` and `sb1p` for categorical.
- **High cardinality with rare levels?** `sketch`, `mba`, `swb`.
- **None of the above:** keep `algorithm = "auto"`.

Points and scaling

```
sc <- obwoe_scale(pdo = 20, score_ref = 600,
                 odds_ref = 50,
                 direction = "higher_is_safer")
s <- obwoe_score(link, sc, round = TRUE)
```

Score = Offset – Factor · η , Factor = PDO / $\ln 2$

The models predict the log-odds of the *event*, so the sign is negative: a high WoE is a risky bin and must score *fewer* points. Two identities pin the scale — a case at `odds_ref` scores exactly `score_ref`, and doubling the good-to-bad odds adds exactly `pdo` points, everywhere.

Tune the binning

```
control.obwoe(
  bin_cutoff = 0.05,           # min share per bin
  max_n_prebins = 20,         # pre-merge ceiling
  convergence_threshold = 1e-6,
  max_iterations = 1000,
  bin_separator = "%;%%",     # joins merged levels
  verbose = FALSE)
```

max_n_prebins is a modelling decision, not a detail. *Pre-binning runs before the optimiser sees the data, so a heavy tail can be smeared into one quantile cell and lost with no warning. Leave the default, and tune it against held-out IV for long-tailed continuous predictors only.*

Call one algorithm directly

```
r <- ob_numerical_ir(german$duration, y,
  min_bins = 3, max_bins = 5,
  bin_cutoff = 0.05, max_n_prebins = 20,
  auto_monotonicity = TRUE)
```

```
r <- ob_categorical_jedi(german$purpose, y,
  bin_separator = "%;%")
```

Every name in the table is exported as `ob_numerical_<name>()`, `ob_categorical_<name>()`, or both. They return the same bin, woe, iv and count vectors that `obwoe()` stores per feature.

The whole pipeline in one call

```
card <- obwoe_scorecard(german,
  target = "default", split = 0.7,
  validation = NULL,           # an out-of-time frame
  exclude = NULL, feature = NULL,
  binning = list(max_bins = 6),
  screening = list(iv_max = 0.5),
  engine = "glm",              # "obwoe" | "glmnet" | list()
  control = control.obwoe_scorecard(),
  file = "scorecard.xlsx", seed = 42)
```

Split, bin, screen by IV and correlation, fit, scale to points, and write the model document. Three properties are enforced rather than assumed:

- Binning is fitted on the **training rows only**, so the hold-out is not inflated by supervised leakage.
- A variable whose WoE coefficient comes out **negative** is dropped and the model refitted — the WoE already carries the direction.
- The generated points SQL reproduces the R card score **exactly**, unseen categories included.

READING THE OBJECT

`card$points` The fixed points table, one row per variable and bin.
`card$screening` `obwoe_select()` plus a stage column: `in_model`, `sign_rejected`, `corr_pruned`, `constant_woe`, `screened_out`.
`card$samples` Per sample: scores, gains table, headline metrics.
`card$stability` PSI for the score and for every variable in the model.
`card$warnings` Everything the pipeline wants a human to read.

```
predict(card, new_data, type = "score")
# "card" | "link" | "prob" | "woe"
```

Redundancy and stability

```
obcorr(df, method = "all", threads = 0)
```

pearson spearman kendall hoeffding distance biweight pbend
robust alternative

```
p <- obwoe_prune(woe_df, ranking = keep,
                 cutoff = 0.7, method = "pearson")
p$keep; p$dropped # variable, correlated_with, correlation
```

Prune in the **WoE space** — that is the space the model actually sees.

```
s <- obwoe_psi(base, compare, n_groups = 10)
s$psi; s$table; s$flag
```

PSI	FLAG	READING
< 0.10	stable	No action
0.10 – 0.25	watch	Monitor the shift
> 0.25	act	Population has moved

*A band populated in one vintage and empty in the other reports **In $\bar{f}$** rather than being smoothed away — a segment that has vanished is exactly what monitoring exists to catch.*

tidymodels

```
rec <- recipe(default ~ ., data = german) |>
  step_obwoe(all_predictors(), outcome = "default",
             algorithm = "auto",
             min_bins = 2, max_bins = 10,
             bin_cutoff = 0.05, output = "woe",
             na_woe = 0)
```

`prep()`, `bake()`, `tidy()` and `required_pkgs()` are all implemented. Four parameters are tunable, each with a dials constructor:

TUNE()	DIALS	RANGE
<code>algorithm</code>	<code>obwoe_algorithm()</code>	28 names
<code>min_bins</code>	<code>obwoe_min_bins()</code>	2 – 5
<code>max_bins</code>	<code>obwoe_max_bins()</code>	5 – 20
<code>bin_cutoff</code>	<code>obwoe_bin_cutoff()</code>	0.01 – 0.10

Deploy to SQL

STEP 4

```
obwoe_sql(m, table = "risk.applications",
  features = keep,
  output = "woe",           # "bin" | "index" | "both"
  style = "select",         # "case" | "cte" | "view"
  dialect = "postgres",
  na_value = 0, explicit_bounds = TRUE,
  quote_identifiers = "auto", file = NULL)
```

ansi postgres mysql mariadb sqlserver oracle spark hive
databricks bigquery snowflake redshift duckdb sqlite

```
CASE WHEN duration IS NULL THEN 0
  WHEN duration <= 7 THEN -1.31218638896617
  WHEN duration > 7 AND duration <= 10 THEN -0.4519851
  ... ELSE 0 END AS duration_woe
```

*Every expression opens with an explicit **IS NULL** branch, because **NULL <= 5** is **NULL** in SQL, not **FALSE**.*